

Lab 3: Expectation Maximisation

Patrick de Kok

October 12, 2012

In this lab assignment, the datasets of both `banana.mat` and `spiral.mat` have been analyzed.

1 Gaussian Distribution

In this exercise, I have trained a classifier, based on a single Gaussian distribution for each class of both datasets.

For `testing`, I use the last `30%` of both datasets. Both datasets have `not been shuffled`. The train and test sets of `banana.mat` and `spiral.mat` have been plotted in `Figure 1` and `Figure 2`.

The training of the classifiers has been done as follows. For both classes C_A and C_B , the mean $\boldsymbol{\mu} = [\boldsymbol{\mu}_A, \boldsymbol{\mu}_B]$ and covariances $\boldsymbol{\Sigma} = [\boldsymbol{\Sigma}_A, \boldsymbol{\Sigma}_B]$ are computed. For this assignment, I have assumed that the frequencies of C_A and C_B are representative of the real distribution. Therefore, $p(C_k)$ is equal to the number of data points of class C_k , divided by the total number of data points. I then compute the most likely class $C_{ML}(\mathbf{y})$ for a data point \mathbf{y} :

$$\begin{aligned} C_{ML}(\mathbf{x}) &= \arg \max_{C_k} p(\mathbf{x}|C_k) \\ &= \arg \max_{C_k} \frac{p(C_k|\mathbf{x})p(C_k)}{\sum_j p(\mathbf{x}|C_j)p(C_j)} \\ &= \arg \max_{C_k} p(C_k|\mathbf{x})p(C_k) \\ &= \arg \max_{C_k} \mathcal{N}(\mathbf{x}; \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)p(C_k) \end{aligned}$$

In MATLAB, the probability density function $\mathcal{N}(\mathbf{x}; \boldsymbol{\mu}, \boldsymbol{\Sigma})$ is computed by `mvnpdf(x, MU, SIGMA)`.

This computation for `banana.mat` results in the confusion matrix and error rate:

$$\begin{aligned} CM(C_{ML}, \text{banana.mat}) &= \begin{bmatrix} 273 & 41 \\ 27 & 259 \end{bmatrix} \\ err(C_{ML}, \text{banana.mat}) &= 0.1133 \end{aligned}$$

Figure 1: A plot of the dataset `banana.mat`. The red “+”’s and magenta “o”’s represent the train and test data points from class A. The blue “x”’s and cyan “*”’s represent the train and test data points from class B.





Figure 2: A plot of the dataset `spiral.mat`. The red “+”’s and magenta “o”’s represent the train and test data points from class A. The blue “×”’s and cyan “*”’s represent the train and test data points from class B.



The corresponding values for `spiral.mat` are:

$$CM(C_{ML}, \text{spiral.mat}) = \begin{bmatrix} 219 & 110 \\ 81 & 190 \end{bmatrix}$$

$$err(C_{ML}, \text{banana.mat}) = 0.3183$$

2 Mixtures of Gaussians

In the second exercise, a classifier will be trained to recognize the classes of the `banana.mat` and `spiral.mat` datasets, based on the EM algorithm.

The iterative loop of the EM algorithm has been implemented already, as well as a visualization module, which is called in every iteration. To make everything work, I have implemented the following three MATLAB functions.

- `MOG = init_mog(X, C)`. Given the data and the desired number of Gaussians, compute the mean `MU`, covariance `SIGMA` and mixing coefficient `PI` for each Gaussian. For N data points, each Gaussian mixture is fitted through $N_k = N/C$ points, and the mean and covariance is computed for this set. `PI` is set to N_k . This is stored in `MOG`, a `C`-by-1 cell array of structures containing only these three keys. For the k th element of `MOG`, `MU`, `SIGMA`

and PI in the software correspond with the mathematical concepts of $\boldsymbol{\mu}_k$, $\boldsymbol{\Sigma}_k$ and π_k .

- `[Q LL] = mog_E_step(X, MOG)`. This implements the expectation step of the EM algorithm. For each data point $\mathbf{x} = \mathbf{X}(\mathbf{n}, :)$ and the k th mixture with mean $\boldsymbol{\mu}_k$, covariance $\boldsymbol{\Sigma}_k$ and mixing coefficient π_k , the responsibility $\pi_k \mathcal{N}(\mathbf{x}; \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$ is divided by the sum of all responsibilities of the k th mixture, and stored in $\mathbf{Q}(\mathbf{n}, k)$, a N -by- C matrix. This corresponds with equation 9.13 of [1]:

$$\begin{aligned} \mathbf{Q}(\mathbf{n}, k) = \gamma(z_k) &\equiv p(z_k = 1 | \mathbf{x}) = \frac{p(z_k = 1)p(\mathbf{x} | z_k = 1)}{\sum_{j=1}^C p(z_j = 1)p(\mathbf{x} | z_j = 1)} \\ &= \frac{\pi_k \mathcal{N}(\mathbf{x}; \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}{\sum_{j=1}^C \pi_j \mathcal{N}(\mathbf{x}; \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)} \end{aligned}$$

However, in my implementation, only one iterator is used; MATLAB can compute the Gaussian distribution for a N -by- D matrix, given a 1-by- D mean and D -by- D covariance matrix. Another speedup is computing the denominator \mathbf{d} together with the numerator, and only dividing each $\mathbf{Q}(\mathbf{n}, k)$ when the complete denominator is known.

LL is the log-likelihood of the dataset under the mixture model. This is computed corresponding to equation 9.14 of [1]:

$$\text{LL} = \ln p(\mathbf{X} | \boldsymbol{\pi}, \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \sum_{n=1}^N \ln \left\{ \sum_{k=1}^C \pi_k \mathcal{N}(\mathbf{x}_n; \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) \right\}$$

Because the innermost sum is the same as \mathbf{d} , the implementation of this formula only takes marginally more time when \mathbf{Q} is already computed.

- `MOG = mog_M_step(X, Q, MOG)`. This function updates the mean MU , covariance SIGMA and mixing coefficient PI in each cell of MOG . For each k th element, this is done according to equations 9.17, 9.19 and 9.22 of [1]:

$$\begin{aligned} \boldsymbol{\mu}_k &= \frac{1}{N_k} \sum_{n=1}^N \gamma(z_{nk}) \mathbf{x}_n \\ \boldsymbol{\Sigma}_k &= \frac{1}{N_k} \sum_{n=1}^N \gamma(z_{nk}) (\mathbf{x}_n - \boldsymbol{\mu}_k)(\mathbf{x}_n - \boldsymbol{\mu}_k)^T \\ \pi_k &= \frac{N_k}{N} \end{aligned}$$

where N is the number of data points in \mathbf{X} , \mathbf{x}_n corresponds with the n th row of \mathbf{X} and N_k is defined by equation 9.18:

$$N_k = \sum_{n=1}^N \gamma(z_{nk})$$

such that $N = \sum_{k=1}^C N_k$.

There are some things that do not follow straightforward from these definitions to the implementation.

Note that $Q(n, k)$ corresponds with $\gamma(z_{nk})$. N_k is implemented as a 1-by- C row vector, generated by summing over all rows of Q .

Because the mean and data points are stored as row vectors, the transpose operator is placed on the first difference vector in the computation of Σ_k :

$$\Sigma_k = \frac{1}{N_k} \sum_{n=1}^N Q(n, k) (\mathbf{x}_n - \boldsymbol{\mu}_k)^T (\mathbf{x}_n - \boldsymbol{\mu}_k)$$

Σ_k is only updated if it is non-singular. In the code, this is defined as $\text{cond}(\Sigma_k) < 10^{10}$. If Σ_k is **singular**, the Gaussian has a really small area (almost no area) with high probabilities, so that **almost no data point will be covered by it**. This is hinted at in the exercise.

Finding the most likely class for the data is done in a similar fashion as for the previous exercise:

$$\begin{aligned} C_{ML}(\mathbf{x}) &= \arg \max_{C_k} p(\mathbf{x}|C_k) \\ &= \arg \max_{C_k} p(C_k) \sum_j p(\mathbf{x}_n|z_j, C_k) p(z_j|C_k) \\ &= \arg \max_{C_k} p(C_k) \sum_j \mathcal{N}(\mathbf{x}_n; \boldsymbol{\mu}_{jk}, \Sigma_{jk}) p(z_j|C_k) \\ &= \arg \max_{C_k} p(C_k) \sum_j \mathcal{N}(\mathbf{x}_n; \boldsymbol{\mu}_{jk}, \Sigma_{jk}) \pi_{jk} \end{aligned}$$

I have tested the test data of both datasets for different models involving 1 through 40 Gaussian mixtures. The error rate of the test set of the **banana.mat** dataset is plotted against the number of Gaussian mixtures in Figure 3. Similar data is plotted in Figure 4 for the **spiral.mat** data set.

For the **banana.mat** dataset, the lowest error rate is 0.0117, which occurred with 10, 18 and 26 Gaussian mixtures. According to Occam's razor, one should always prefer the simplest hypothesis, I would advise using **$C = 10$** .

I find the same number of Gaussians for the **spiral.mat** dataset. The lowest error rate is 0.0017, for $C \in \{10, 16\}$. With the same argumentation, I advise **10 Gaussians**.

Figure 3: The number of Gaussians plotted against the error rate for the `banana.mat` dataset.


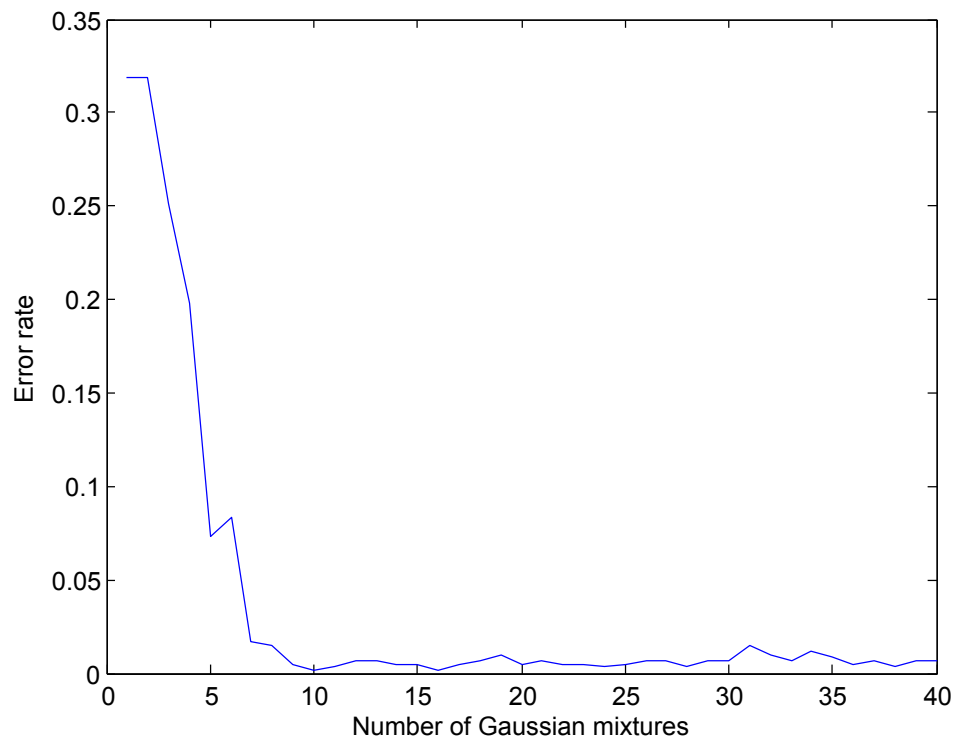


Figure 4: The number of Gaussians plotted against the error rate for the `spiral.mat` dataset.



3 Log-Probabilities

Because of **underflow errors** occurring when multiplying floating point numbers, one might like to work with **logaritms of probabilities**, instead of directly with probabilities. Multiplication of two logarithms $\ln(a)\ln(b)$ is the same as the logarithm of a sum $\ln(a + b)$. To do this, we need to compute that new value.

For two N -dimensional vectors x, y , determine per dimension the maximum and minimum element, and store these two vector as *high* and *low*. Then the logarithm of the sum is computed through **$high + \ln(1 + \exp(low - high))$** ¹.

The function is implemented in such a way that it **computes the sum by adding the i th row by the intermediate result**. By doing it in this way, `logsumexp` accepts a general matrix as input, and computes the sum along the rows. The implementation computes the correct sum. For example **`logsumexp([-1000, -1001]) = -999.6867`**.

The expectation step should be slightly modified. Where we used before

$$Q(n, k) = \frac{\pi_k \text{mvnpdf}(\mathbf{x}, \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}{\sum_{j=1}^C \pi_j \text{mvnpdf}(\mathbf{x}, \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}$$

we now compute:

$$Q(n, k) = \exp \left(\ln \pi_k \text{lmvnpdf}(\mathbf{x}, \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) - \sum_{j=1}^C \pi_j \text{lmvnpdf}(\mathbf{x}, \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) \right)$$

This did not change the results for both datasets, both for the best number of Gaussian mixtures for classification, as the error rates.

References

- [1] Christopher M. Bishop. *Pattern Recognition and Machine Learning (Information Science and Statistics)*. Springer-Verlag New York, Inc., Secaucus, NJ, USA, 2006.

¹Note that in the implementation, the scalar 1 must be replaced by a N -dimensional vector with all 1s.